




MATH-500 (n=500)

AIME24 (n=960)





AIME25 (n=960)



Olympiad (n=675)

AMC23 (n=1280)

Minerva (n=272)



T: Token count | C: Router confidence | M: Leading-expert margin | S: Selected probability mass | B: Selection-boundary margin | G: Selected/rest mean gap
H: Hidden-state norm | D: Hidden-state step distance | R: Hidden/router geometry | W: Top-1 switch rate | O: Adjacent selected-set overlap