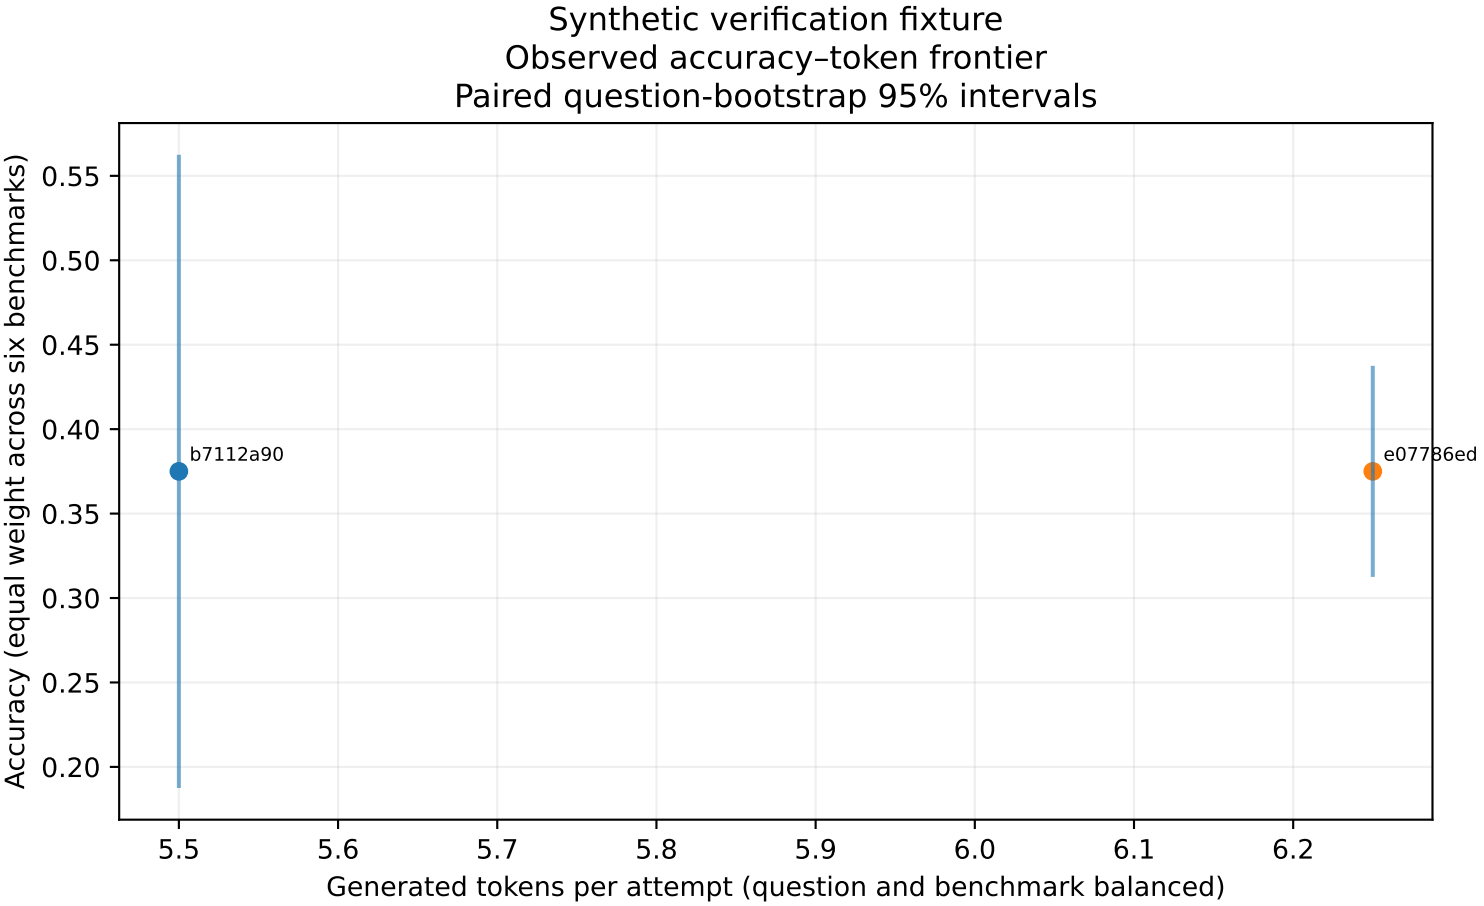
Synthetic verification fixture
Observed accuracy–token frontier
Paired question-bootstrap 95% intervals
Accuracy (equal weight across six benchmarks)
0.55
0.50
0.45
0.40
0.35
0.30
0.25
0.20
b7112a90
e07786ed
5.5
5.6
5.7
5.8
5.9
6.0
6.1
6.2
Generated tokens per attempt (question and benchmark balanced)